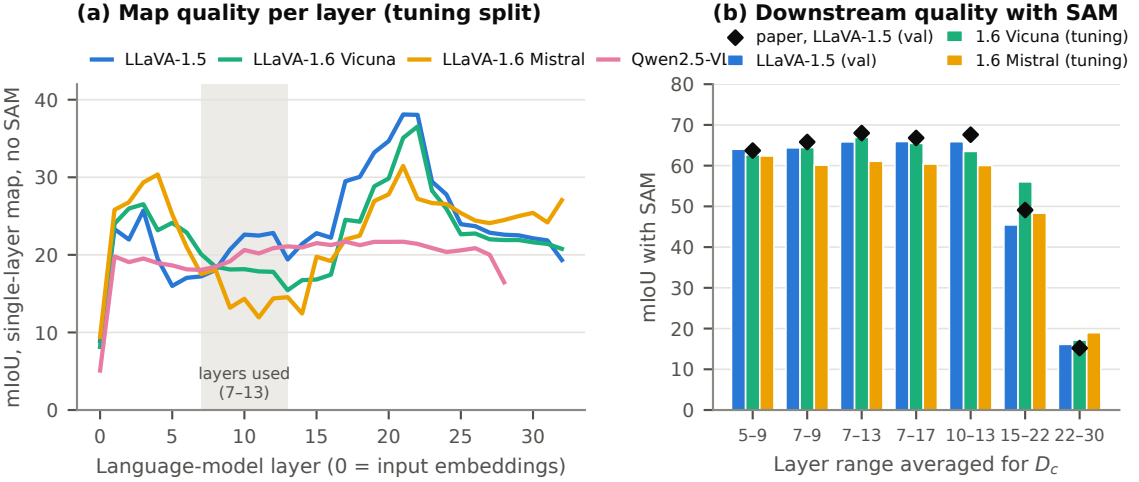
(a) Map quality per layer (tuning split)
LLaVA-1.5
LLaVA-1.6 Vicuna
LLaVA-1.6 Mistral
Qwen2.5-VL
mIoU, single-layer map, no SAM
layers used
(7–13)
Language-model layer (0 = input embeddings)
(b) Downstream quality with SAM
paper, LLaVA-1.5 (val)
LLaVA-1.5 (val)
1.6 Vicuna (tuning)
1.6 Mistral (tuning)
mIoU with SAM
5–9
7–9
7–13
7–17
10–13
15–22
22–30
Layer range averaged for D_c